

atharva-bhide-assignment-4

April 10, 2024

1 Import libraries

```
[3]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.datasets import load_boston
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, mean_absolute_error
from sklearn.preprocessing import StandardScaler
import warnings
warnings.filterwarnings("ignore")
%matplotlib inline
```

```
[4]: boston = load_boston()
boston.keys()
```

```
[4]: dict_keys(['data', 'target', 'feature_names', 'DESCR', 'filename',
'data_module'])
```

```
[5]: x = pd.DataFrame(boston.data, columns=boston.feature_names)
y = pd.DataFrame(boston.target, columns=['MEDV'])
```

```
[6]: x.head()
```

```
[6]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	

	PTRATIO	B	LSTAT
0	15.3	396.90	4.98
1	17.8	396.90	9.14
2	17.8	392.83	4.03

```
3      18.7  394.63   2.94
4      18.7  396.90   5.33
```

```
[7]: x.shape, y.shape
```

```
[7]: ((506, 13), (506, 1))
```

2 Basic stats

```
[8]: x.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 13 columns):
#   Column      Non-Null Count  Dtype
---  -
0   CRIM        506 non-null    float64
1   ZN          506 non-null    float64
2   INDUS       506 non-null    float64
3   CHAS        506 non-null    float64
4   NOX         506 non-null    float64
5   RM          506 non-null    float64
6   AGE         506 non-null    float64
7   DIS         506 non-null    float64
8   RAD         506 non-null    float64
9   TAX         506 non-null    float64
10  PTRATIO     506 non-null    float64
11  B           506 non-null    float64
12  LSTAT       506 non-null    float64
dtypes: float64(13)
memory usage: 51.5 KB
```

```
[9]: x.describe()
```

```
[9]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM \
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000

	AGE	DIS	RAD	TAX	PTRATIO	B \
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000

mean	68.574901	3.795043	9.549407	408.237154	18.455534	356.674032
std	28.148861	2.105710	8.707259	168.537116	2.164946	91.294864
min	2.900000	1.129600	1.000000	187.000000	12.600000	0.320000
25%	45.025000	2.100175	4.000000	279.000000	17.400000	375.377500
50%	77.500000	3.207450	5.000000	330.000000	19.050000	391.440000
75%	94.075000	5.188425	24.000000	666.000000	20.200000	396.225000
max	100.000000	12.126500	24.000000	711.000000	22.000000	396.900000

```

LSTAT
count    506.000000
mean      12.653063
std        7.141062
min        1.730000
25%        6.950000
50%       11.360000
75%       16.955000
max       37.970000

```

```
[10]: y.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 1 columns):
 #   Column  Non-Null Count  Dtype  
---  -
 0   MEDV    506 non-null     float64
dtypes: float64(1)
memory usage: 4.1 KB

```

```
[11]: y.describe()
```

```

[11]:          MEDV
count    506.000000
mean      22.532806
std        9.197104
min         5.000000
25%       17.025000
50%       21.200000
75%       25.000000
max       50.000000

```

```
[12]: x.isnull().sum()
```

```

[12]: CRIM      0
      ZN        0
      INDUS    0
      CHAS     0

```

```

NOX      0
RM       0
AGE      0
DIS      0
RAD      0
TAX      0
PTRATIO  0
B        0
LSTAT    0
dtype: int64

```

```
[13]: y.isnull().sum()
```

```
[13]: MEDV      0
dtype: int64
```


```
[14]: df = x
df["target"] = y
df.head()
```

```
[14]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	

	PTRATIO	B	LSTAT	target
0	15.3	396.90	4.98	24.0
1	17.8	396.90	9.14	21.6
2	17.8	392.83	4.03	34.7
3	18.7	394.63	2.94	33.4
4	18.7	396.90	5.33	36.2


```
[15]: plt.figure(figsize=(15,10))
sns.heatmap(df.corr(), annot=True)
plt.show()
```



2.0.1 Considering only 'RM' and 'LSTAT' by considering correlation and multicollinearity of other features

```
[51]: df = df[['RM', 'LSTAT', 'target']]
```

```
[52]: sns.pairplot(df)
plt.show()
```



```
[63]: x = df[['RM', 'LSTAT']]
      y = df['target']
```

3 Scale the data

```
[64]: scaler = StandardScaler()
```

```
[65]: x = scaler.fit_transform(x)
```

4 Split the data

```
[67]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3, shuffle=True)
```

```
[68]: x_train.shape, x_test.shape, y_train.shape, y_test.shape
```

```
[68]: ((354, 2), (152, 2), (354,), (152,))
```

5 Linear Regression Modelling

```
[69]: model = LinearRegression(n_jobs=-1)
```

```
[70]: model.fit(x_train, y_train)
```

```
[70]: LinearRegression(n_jobs=-1)
```

6 Make predictions

```
[71]: y_pred = model.predict(x_test)
```

```
[72]: mean_absolute_error(y_test, y_pred)
```

```
[72]: 3.701010266760501
```

```
[73]: mean_squared_error(y_test, y_pred)
```

```
[73]: 30.5001478179898
```

```
[74]: sns.regplot(y_test, y_pred, color='red')  
plt.show()
```

